

Tuesday 31st October 2023


Big Data

Data Mining, warehouse, viz, nosql
applied to health data


Agenda




Data streamin
introduction



Request/Response
Vs
Event-Driven



Kafka



Data streaming

What ?

Real-Time processing

Process data as it arrives

Continuous


Unending stream,
from multiple sources

Timeliness

Quick /
instant insights from
incoming data

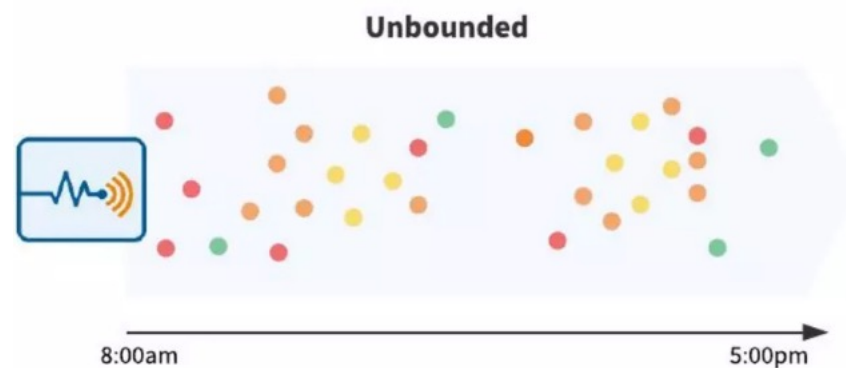
Batch vs Streaming

Batch processing



Processing on accumulated data on a specific bounded period. Ideal for huge n-time-critical data

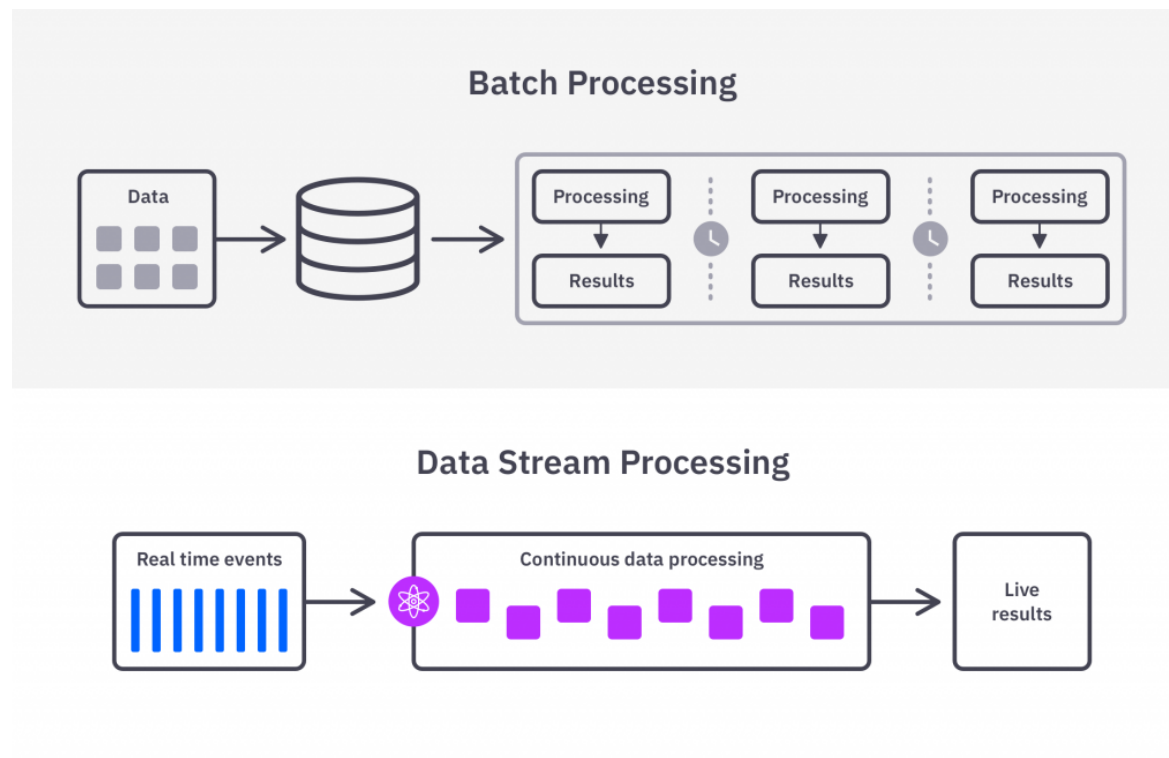
Data Streaming (Real-time)



Continuous and instantaneous processing of cross-devices data as it arrives

Batch vs Streaming

Source: <https://addepto.com/blog/stream-data-model-and-architecture/>



Characteristics of Data Streams

Size

Small in size (in KBs)

Volume

Unbounded =>
very large
volume because
of accumulation

Examples

lot sensors, log
files, financial
markets, apps
clicks

Velocity

Fast Changing,
requires fast
reponse

Variety

Unstructured or
semi-structured

Challenges

Scalability

Handling Large
volumes at high
velocity

Veracity

Ensuring data
quality and
consistency

Interagtion

Integrate
existing systems
and infras

Examples


Some concrete use-cases



Finance

Fraud Detection – Analyzing patterns and anomalies to detect fraudulent activities, unauthorized transactions


Financial transactions – Analyze market trade and execute trades in real-time



Health

Patients Monitoring – Stream patient's vital signs thanks to IoT sensors, enabling immediate intervention

Predictive Analytics for disease risk -



Media

Content Personalization

Real time audience analytics for media streaming platforms

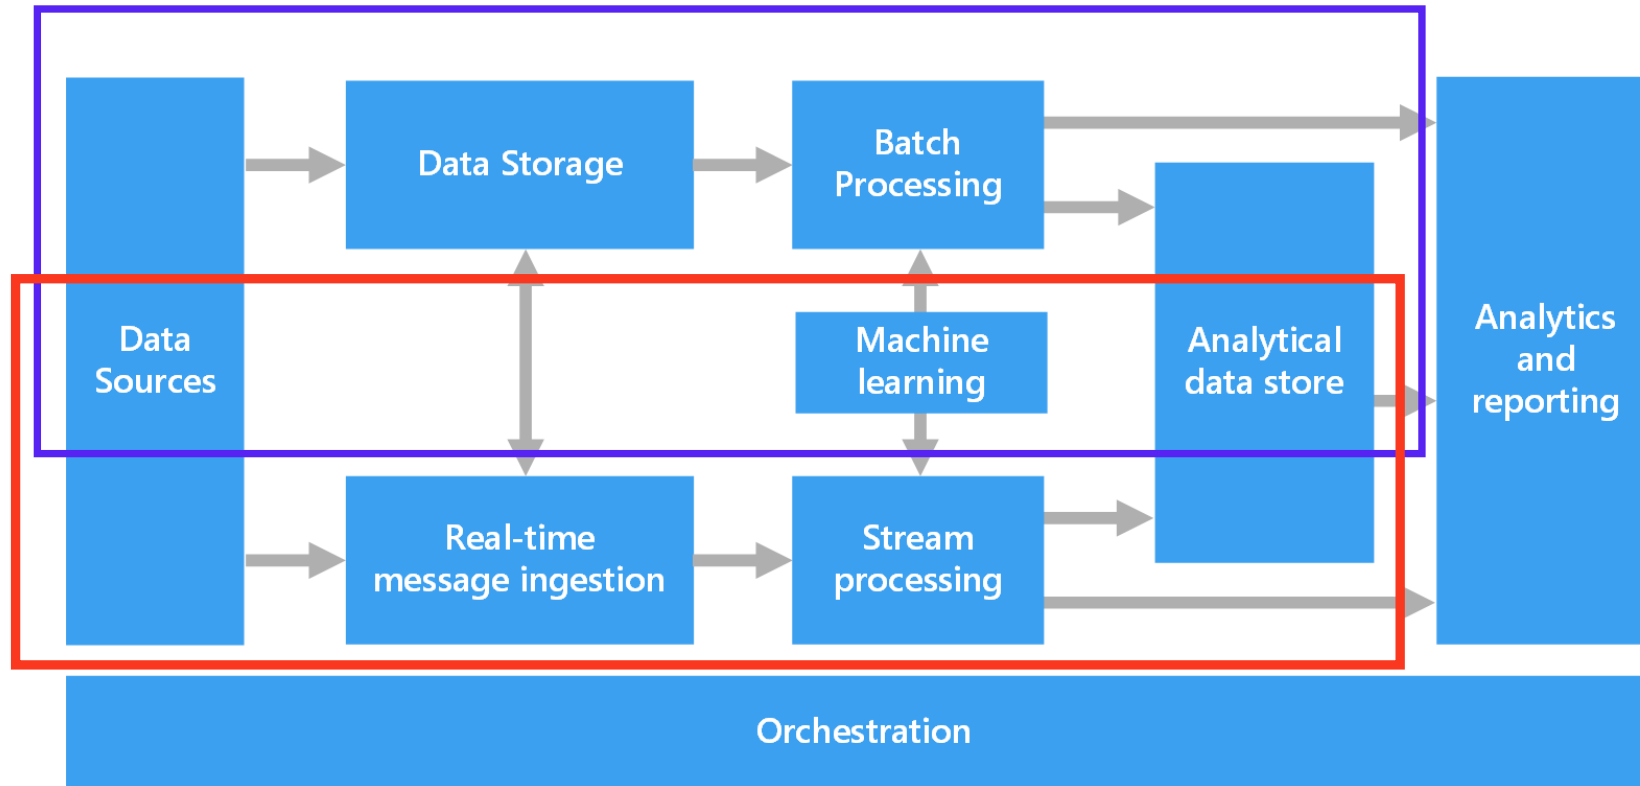


Retail

Dynamic Pricing and inventory management


Enhancing customer experience through real-time insights

Reminder : Big Data Architecture




How do you stream data ?

From sources to datalake ?



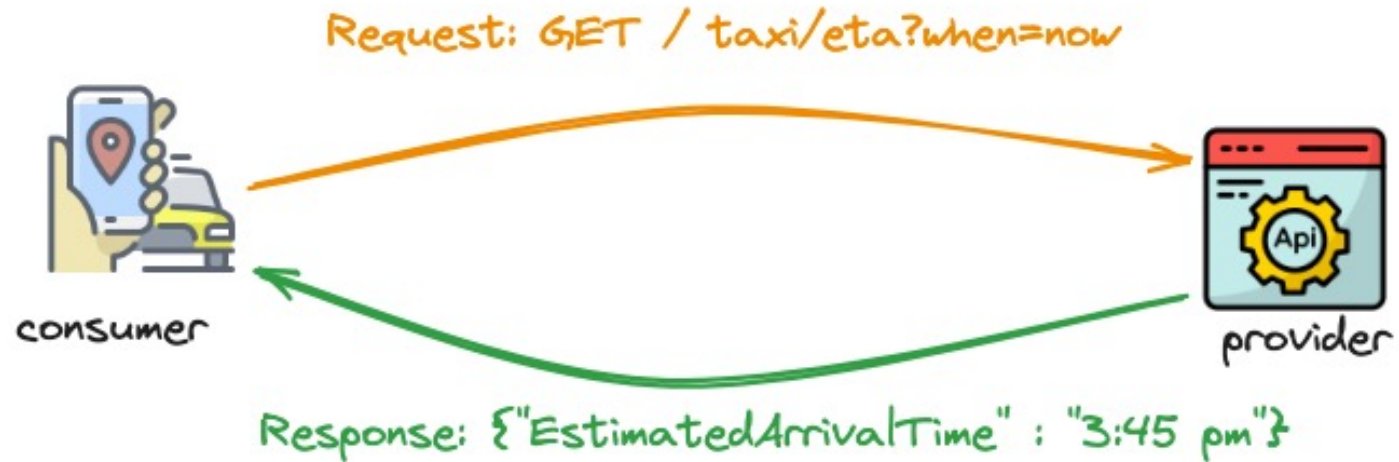
First approach : RESTful API

From sources to datalake ?



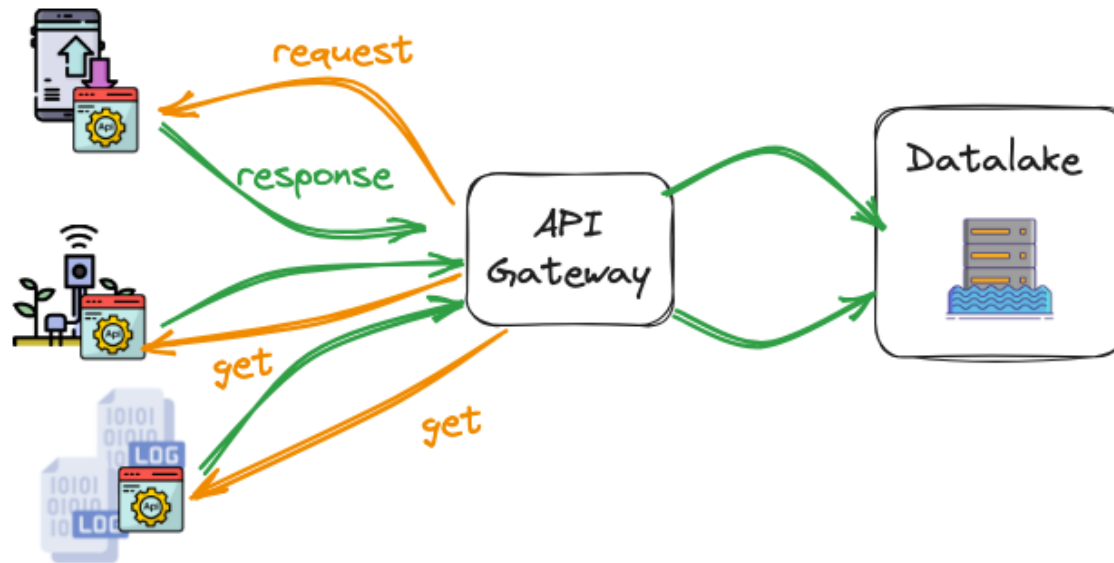
First approach : RESTful API

From sources to datalake ?



First approach : RESTful API

From sources to datalake ?



Pros :


- Simple : Largely used and understood in the industry
- Interoperability : Easy integration for different systems
- Same HTTP norms everywhere

Cons :

- Active and continuous demand
- Latency : periodical « GET » may induce latency
- Scalability : difficult to maintain as data grows


First approach : RESTful API

Not really "real-time"




Second approach : Event-driven

Be smarter when asking for your ride




Second approach : Event-driven

Subscription to events



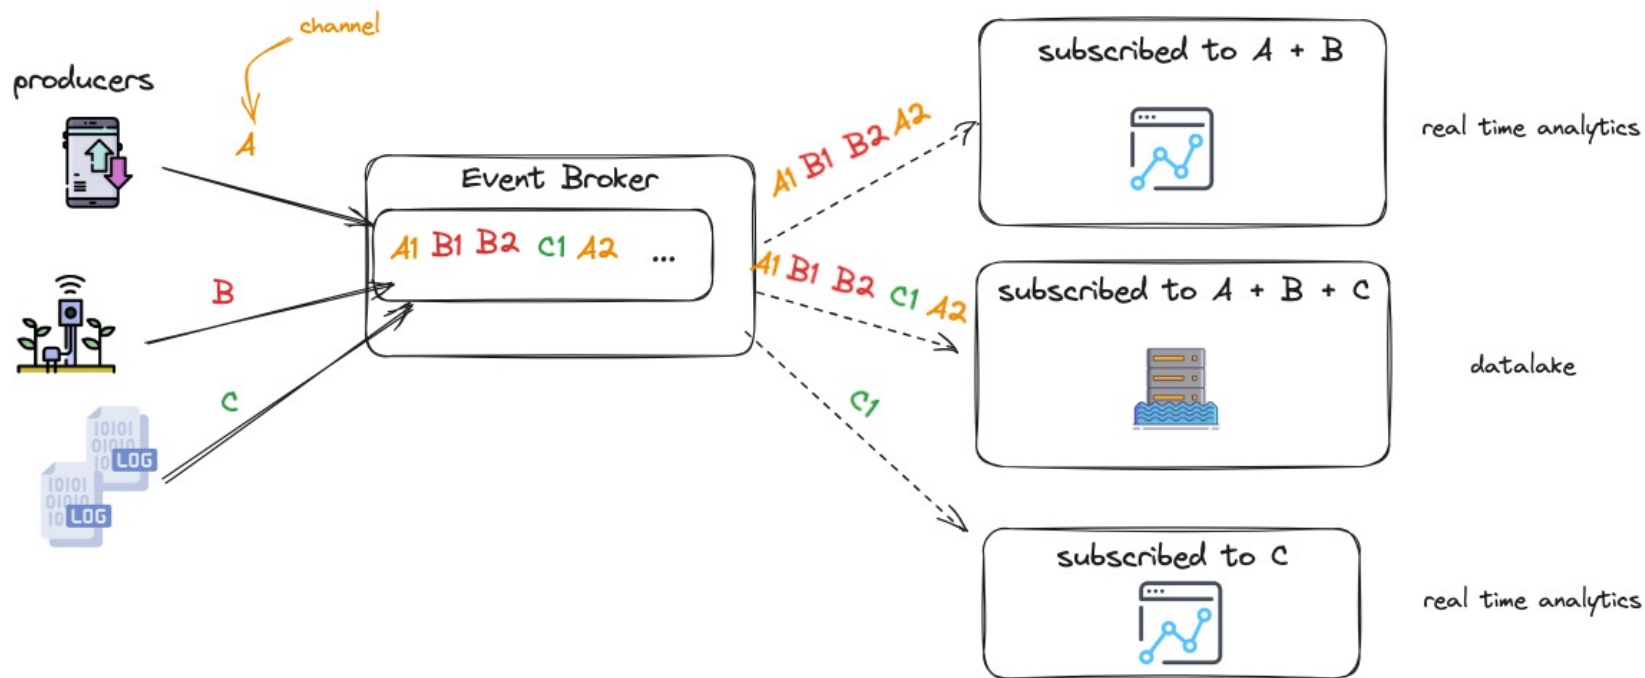
Second approach : Event-driven

Subscription to events




Second approach : Event-driven

Event-driven real-time architecture

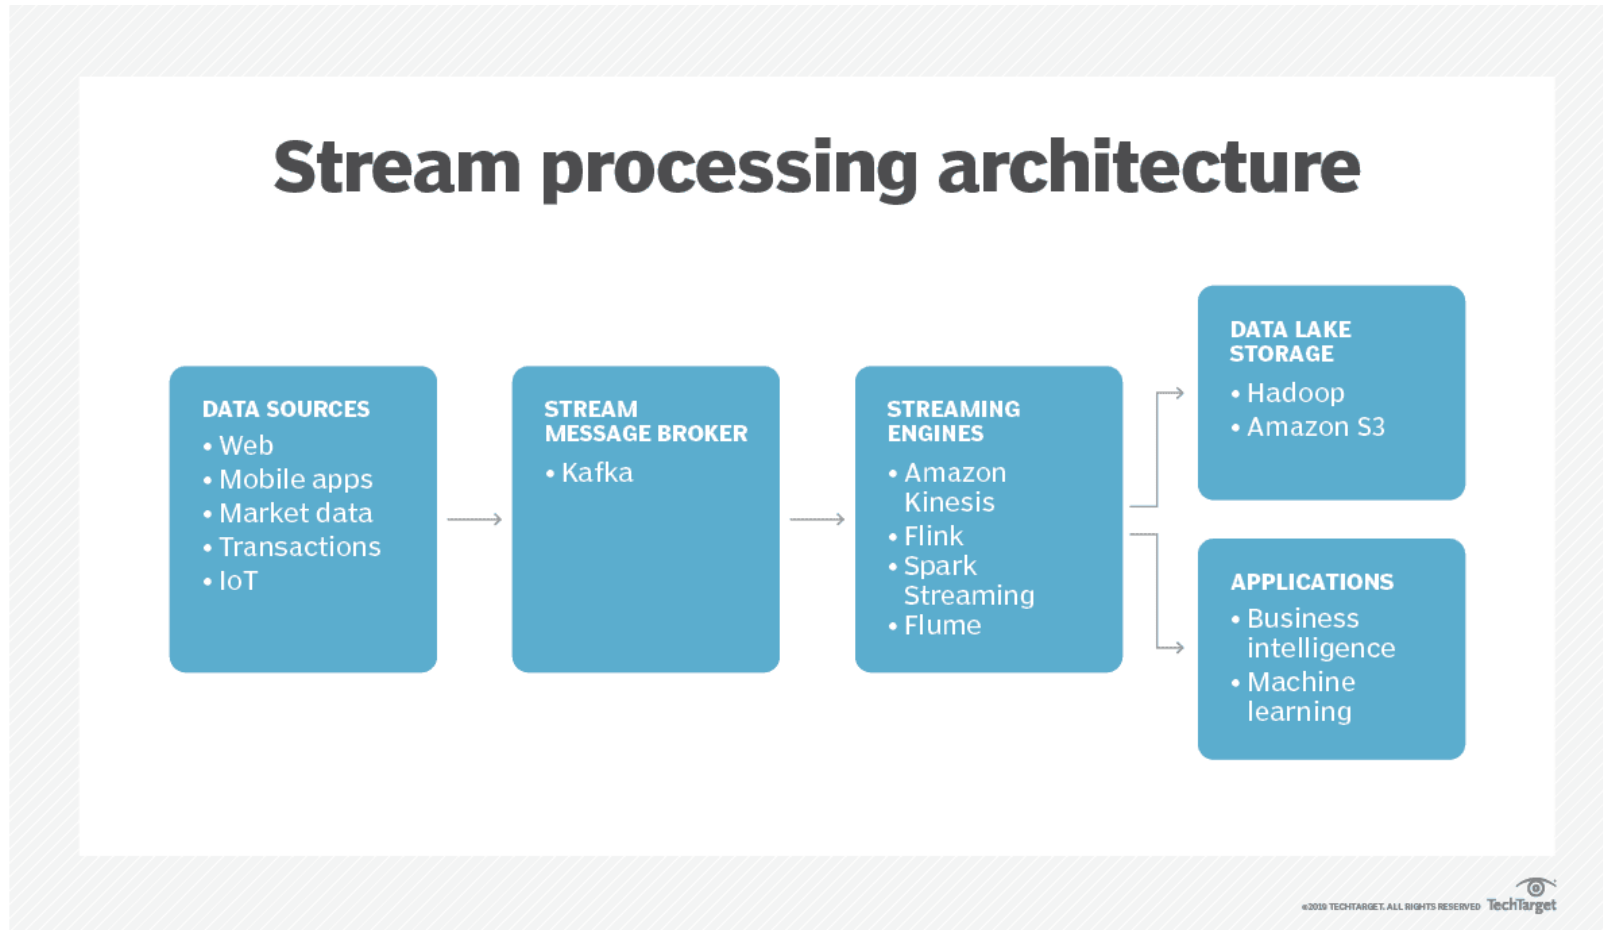


- Asynchronous
- Real-time automatic events **detection**
- Producers and consumers are **independent**
- **Scalability** : easy to add more producer or consumer
- All services are **independants**

Data Streaming Architecture



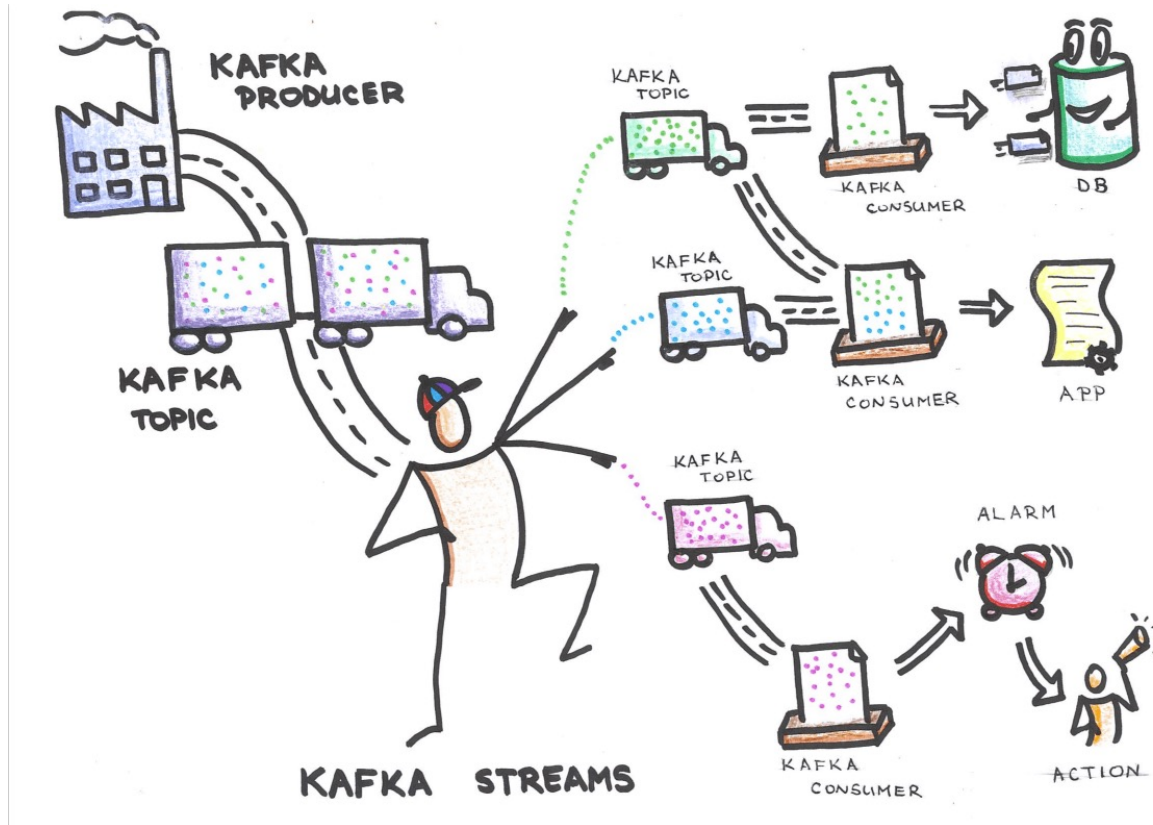
Data Streaming Architecture



Apache Kafka

Introduction to Apache Kafka

Apache Kafka: Real-Time Data Streaming Platform

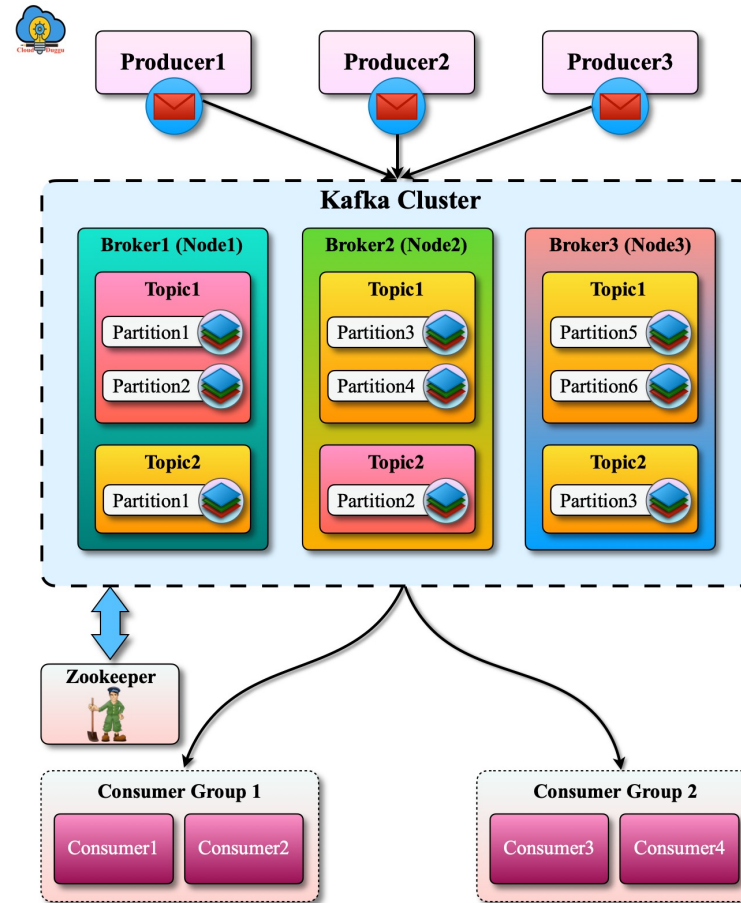


Source : <https://croz.net/news/apache-kafka/>

- Asynchronous
- Real-time automatic events **detection**
- Producers and consumers are **independent**
- **Scalability** : easy to add more producer or consumer
- All services are **independants**

Kafka Architecture Overview

Understanding Kafka's Core Components



- **Producers** : Services that **publish** data to Kafka topics
- **Consumers** : Services or apps that **subscribe** to topics and process data
- **Brokers** : Kafka servers that store data and serve clients
- **Zookeeper** : Manages Kafka cluster metadata and brokers

Kafka's key features

The power of Kafka for Real-Time Data Operations



- **Scalable** in all dimensions, including event producers, processors, consumers, and connectors
- **Durability** and **Reliability**: data is replicated when stored on disk, ensuring data persistence and reliability
- **>1M Throughput** for both publishing and subscribing
- **Real-Time Processing** + other technologies integration
- **Fault Tolerance**
- **Distributed Architecture**

Kafka

Usecases

